

Series 9: Neural Networks

Computational Statistics, FS 2022

May 2022

1 Measuring Distances between Distributions

Recall from Exercise Class 10 that the Kullback-Leiber Divergence gives us a way to quantify the difference between two probability distributions. For continuous probability densities $P(x)$ and $Q(x)$, it is defined as

$$\text{KL}(P, Q) = \int P(x) \log \left(\frac{P(x)}{Q(x)} \right) dx. \quad (1)$$

1. Show that the Kullback-Leibler Divergence satisfies $\text{KL}(P, Q) \geq 0$, where equality happens if and only if $P(x) = Q(x)$ for all x . You may assume that P and Q are continuous.

2. Is it true that $\text{KL}(P, Q) = \text{KL}(Q, P)$?

3. Show that the Kullback-Leibler Divergence between multivariate Gaussian distributions $P = \mathcal{N}(\boldsymbol{\mu}_P, \boldsymbol{\Sigma}_P)$ and $Q = \mathcal{N}(\boldsymbol{\mu}_Q, \boldsymbol{\Sigma}_Q)$ is

$$\text{KL}(P, Q) = \frac{1}{2} \left\{ \log \frac{\det \boldsymbol{\Sigma}_Q}{\det \boldsymbol{\Sigma}_P} - d + \text{Tr}(\boldsymbol{\Sigma}_Q^{-1} \boldsymbol{\Sigma}_P) + (\boldsymbol{\mu}_P - \boldsymbol{\mu}_Q)^\top \boldsymbol{\Sigma}_Q^{-1} (\boldsymbol{\mu}_P - \boldsymbol{\mu}_Q) \right\}, \quad (2)$$

where figuring out the meaning of d is part of the exercise.

The equality (2) plays an important role in implementing a variational auto-encoder (VAE).

2 Deriving the Auto-Encoding Variational Bayes

In this exercise, we derive the auto-encoding variational Bayes algorithm by following the original paper by Kingma and Welling (see references below). The reason we derive this algorithm is that the variational auto-encoder is a special case (see Exercise 3 below).

The framework is a probabilistic model with latent variables. Specifically, we model the input vector $\mathbf{x} \in \mathbb{R}^p$ and latent variables $\mathbf{z} \in \mathbb{R}^d$ with a family of joint distributions

$$p_{\boldsymbol{\theta}}(\mathbf{x}, \mathbf{z}) = p_{\boldsymbol{\theta}}(\mathbf{z}) p_{\boldsymbol{\theta}}(\mathbf{x} | \mathbf{z}) \quad (3)$$

indexed by $\boldsymbol{\theta}$.

We think of the latent variables \mathbf{z} as a lower-dimensional “code” representing \mathbf{x} . The conditional probability $p_{\theta}(\mathbf{x}|\mathbf{z})$ represents *decoding* \mathbf{x} from the code \mathbf{z} . Conversely, the reverse conditional $p_{\theta}(\mathbf{z}|\mathbf{x})$ specifies the *encoding* distribution. To fit our probability model, we approximate the true encoding distribution $p_{\theta}(\mathbf{z}|\mathbf{x})$ by a simpler approximate model $q_{\phi}(\mathbf{z}|\mathbf{x})$. In a variational auto-encoder, this simpler encoding model is a Gaussian distribution, but we stick with the general framework in this derivation.

1. Assume that n training examples $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(n)} \in \mathbb{R}^p$ are sampled i.i.d. Show that the marginal log likelihood $\log p_{\theta}(\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(n)})$ equals

$$\sum_i^n \log p_{\theta}(\mathbf{x}^{(i)}) = \sum_{i=1}^n \log \left(\frac{q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)}) p_{\theta}(\mathbf{z}|\mathbf{x}^{(i)}) p_{\theta}(\mathbf{x}^{(i)})}{p_{\theta}(\mathbf{z}|\mathbf{x}^{(i)}) q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})} \right). \quad (4)$$

2. Focus on the marginal log-likelihood $\log p_{\theta}(\mathbf{x}^{(i)})$ for a single training example $\mathbf{x}^{(i)}$. By rewriting the expression inside the log as in Equation (4), show that

$$\log p_{\theta}(\mathbf{x}^{(i)}) = \text{KL}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)}), p_{\theta}(\mathbf{z}|\mathbf{x}^{(i)})) + \mathcal{L}(\theta, \phi; \mathbf{x}^{(i)}) \quad (5)$$

where

$$\mathcal{L}(\theta, \phi; \mathbf{x}^{(i)}) = \sum_{\mathbf{z}} q_{\phi}(\mathbf{z}|\mathbf{x}) \log \left(\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q_{\phi}(\mathbf{z}|\mathbf{x})} \right). \quad (6)$$

Note that since $\text{KL}(\cdot, \cdot) \geq 0$, the term $\mathcal{L}(\theta, \phi; \mathbf{x}^{(i)})$ is a lower bound on the marginal log-likelihood. This is a specific instance of the *evidence lower bound* (see Wikipedia).

Hint: As long as the log in Equation (4) does not depend on \mathbf{z} , you can take its expectation with respect to any distribution over \mathbf{z} and its value will not change. You may take a discrete expectation, *i.e.* $\sum_{\mathbf{z}} p(\mathbf{z})f(\mathbf{z})$ rather than $\int p(\mathbf{z})f(\mathbf{z})d\mathbf{z}$.

3. To fit this model, we maximize the log likelihood by maximizing its lower bound $\mathcal{L}(\theta, \phi; \mathbf{x}^{(i)})$. In other words, $-\mathcal{L}(\theta, \phi; \mathbf{x}^{(i)})$ is our loss function for training. Rewrite the right hand side of Equation (6) to conclude that

$$\mathcal{L}(\theta, \phi; \mathbf{x}^{(i)}) = -\text{KL}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)}), p_{\theta}(\mathbf{z})) + \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})}[\log p_{\theta}(\mathbf{x}^{(i)}|\mathbf{z})]. \quad (7)$$

Interpret the two terms on the right hand side as follows. The KL term regularizes our encoding distribution towards a prior distribution over the latent variables. The expectation term measures the reconstruction error: it assesses how closely the decoder’s output resembles the encoder’s input.

3 Deriving the Variational Auto-Encoder

The variational auto-encoder (VAE) is a special case of the auto-encoding variational Bayes model derived in Exercise 2 above.

In a variational auto-encoder, optimizing for ϕ in $q_\phi(\mathbf{z}|\mathbf{x})$ means fitting an encoder neural network that computes a Gaussian mean vector and covariance matrix as a function of \mathbf{x} , to specify the encoding distribution. The parameter ϕ represents the weights of this neural network.

Similarly, optimizing for θ in $p_\theta(\mathbf{x}|\mathbf{z})$ means fitting the decoder neural network. We take the decoding distribution $p_\theta(\mathbf{x}|\mathbf{z})$ to be Gaussian with mean vector $\mathbf{x}'(\mathbf{z})$ depending on the latent representation \mathbf{z} and unit covariance matrix. The decoder neural network computes $\mathbf{x}'(\mathbf{z})$ from \mathbf{z} . The parameter θ represents the weights of this neural network.

More formally, a VAE models the encoding and decoding distributions as

$$\begin{aligned} q_\phi(\mathbf{z}|\mathbf{x}) &\sim \mathcal{N}(\boldsymbol{\mu}(\mathbf{x}), \text{diag}(\boldsymbol{\sigma}^2(\mathbf{x}))) \\ p_\theta(\mathbf{x}|\mathbf{z}) &\sim \mathcal{N}(\mathbf{x}'(\mathbf{z}), \mathbf{I}), \end{aligned}$$

where $\boldsymbol{\sigma}^2(\mathbf{x})$ is a vector of variances. In addition, we choose a Gaussian prior $p_\theta(\mathbf{z}) = \mathcal{N}(\mathbf{0}, \mathbf{I})$ over the latent representation. The output of the decoder neural network is the vector \mathbf{x}' . After training a VAE, its main purpose lies in computing lower-dimensional latent representations for input data points. To compute a latent representation \mathbf{z} for an input point \mathbf{x} , we sample from the encoding distribution $q_\phi(\mathbf{z}|\mathbf{x}) \sim \mathcal{N}(\boldsymbol{\mu}(\mathbf{x}), \text{diag}(\boldsymbol{\sigma}^2(\mathbf{x})))$.

In this exercise, we build on the results of Exercise 2 to derive the loss function for training a VAE.

1. Use the results of Exercise 1 to express the KL term in Equation (7) as

$$\text{KL}(\phi) = \frac{1}{2} \left[\left(\sum_{j=1}^d \sigma_j^2(\mathbf{x}^{(i)}) - \log \sigma_j^2(\mathbf{x}^{(i)}) \right) - d + \|\boldsymbol{\mu}(\mathbf{x}^{(i)})\|^2 \right], \quad (8)$$

where ϕ specifies the encoder neural network that maps \mathbf{x} to $\boldsymbol{\mu}(\mathbf{x})$ and $\boldsymbol{\sigma}^2(\mathbf{x})$.

Note: There is no dependence on θ (representing the decoder neural network) because our prior distribution over the latent representation, $p_\theta(\mathbf{z})$, is simply $\mathcal{N}(\mathbf{0}, \mathbf{I})$.

2. Express the reconstruction error (expectation term in Equation (7)) for training example $\mathbf{x}^{(i)}$ as

$$-\frac{1}{2} \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x}^{(i)})} [\|\mathbf{x}^{(i)} - \mathbf{x}'(\mathbf{z})\|^2] + \text{constant term}. \quad (9)$$

3. Conclude that maximizing the lower bound $\mathcal{L}(\theta, \phi; \mathbf{x}^{(i)})$ of the log likelihood is equivalent to minimizing the loss function

$$\ell(\theta, \phi) = \left(\sum_{j=1}^d \sigma_j^2(\mathbf{x}^{(i)}) - \log \sigma_j^2(\mathbf{x}^{(i)}) \right) + \|\boldsymbol{\mu}(\mathbf{x}^{(i)})\|^2 + \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x}^{(i)})} [\|\mathbf{x}^{(i)} - \mathbf{x}'(\mathbf{z})\|^2],$$

where the parameters ϕ and θ specify the mappings $\mathbf{x} \mapsto \boldsymbol{\mu}(\mathbf{x}), \boldsymbol{\sigma}^2(\mathbf{x})$ and $\mathbf{z} \mapsto \mathbf{x}'(\mathbf{z})$ in the form of neural networks (encoder and decoder). Optimizing for θ, ϕ corresponds to optimizing the weights of these two neural networks.

4. (Bonus) It is not obvious how to compute the gradient with respect to ϕ of the expectation term in our loss function. Read the paper by Kingma and Welling to familiarize yourself with their *reparametrization trick*. You are then ready to train a variational auto-encoder with gradient descent.

4 References

- [1] D. P. Kingma and M. Welling. Auto-Encoding Variational Bayes. In *2nd International Conference on Learning Representations, ICLR 2014*, 2014.
→ access here: <https://arxiv.org/abs/1312.6114>